

SentienGuard GTM Playbook V2.3

Autonomous AIOps and controlled infrastructure remediation

V2.3 Research Edition

Evidence-led strategy. No revenue forecast. Internal validation required where noted.

Executive understanding

SentienGuard's strongest distinction is not monitoring. It is the remediation layer after monitoring: Detect -> Diagnose -> Execute -> Verify -> Prove. It can start read-only, move to human approval and then automate proven playbooks while coexisting with Datadog, PagerDuty, Prometheus, Grafana, AWS, Azure and GCP.

1. Product-led commercial thesis

- Use Product-Qualified Account signals, not blanket outreach to every free user.
- Observe -> Approve -> Automate is both a technical trust ladder and a sales progression.
- Use SPICED to quantify toil and critical events; MEDDPICC for enterprise/fleet procurement; Challenger to reframe 'we already have monitoring' into 'monitoring detects, remediation closes the loop'.

2. Pricing and adoption signals

Tier	Current public positioning	Commercial signal
Individual	\$0 for first 3 nodes	Experiment / technical validation
Team	\$5/node monthly or \$4/node annual; unlimited nodes	Team adoption and broader workload coverage
Fleet / MSP	Custom	Multi-org, enterprise or managed-service motion

PQA signal	Why it matters	Sales action
Node growth	Breadth of infrastructure coverage is increasing	Ask what remains manual across the growing estate
Production environment added	Trust and business criticality increased	Offer controlled-remediation review
Multiple users / teams	Experiment becoming organizational adoption	Map Platform/SRE/Security buying committee
Custom playbooks / API use	Operational embedding	Discuss fleet governance and enterprise requirements
Security / procurement questions	Commercial intent signal	Move to enterprise qualification

3. Trigger log

Trigger	Verify	Hypothesis
Kubernetes / cloud footprint growth	Platform/SRE hiring, new regions, cloud expansion	Incident complexity and repetitive operations may be increasing
Reliability / on-call pressure	SRE hiring, incident program, public reliability initiative	Senior engineers may be spending time on known remediation
Compliance / change-control requirement	SOC 2/HIPAA/financial controls, audit program	Evidence and production authority may need stronger governance

4. Competitive intelligence

Competitor / status quo	Strength	SentienGuard wedge	Where not to compete
PagerDuty	Incident response ecosystem, orchestration, enterprise adoption	Remediation-first closed loop with verification/rollback where validated	Do not claim PagerDuty cannot automate
BigPanda	Event correlation and incident intelligence	Execute/verify after correlation; routine remediation layer	If pain is primarily alert correlation, lead cautiously
Datadog / Dynatrace / Grafana etc.	Deep observability and telemetry	Coexist, do not replace; sell resolution layer	Avoid anti-observability positioning
Runbooks/scripts/Ansible/Rundeck	Known automation and operator control	AI-assisted playbook selection, gated authority, verification, rollback, evidence	If customer only needs deterministic manual execution, do not oversell AI

5. Economic impact model

Metric	Customer inputs	Calculation
Incident toil	Incidents/month x engineers x hours/incident x loaded hourly cost	Annualize x 12
Repetitive escalation	Repeatable incidents x senior engineer hours x loaded rate	Annualize x 12
Downtime	Verified downtime hours x customer-provided business loss/hour	Only use customer data
Roadmap opportunity cost	Recovered engineer hours	Keep qualitative unless customer has approved value assumptions

Vendor-claim discipline

The site currently cites 87% routine incidents resolved autonomously in its validation dataset and sub-90-second typical autonomous resolution. Preserve the dataset caveat. Do not present these as universal production outcomes.

6. PLG -> Enterprise motion

Stage	Goal	Exit trigger
Free / Observe	Technical fit, zero autonomous changes	Agents connected, useful diagnosis, no security blocker
Approve	Human-authorized remediation	Known incident classes repeatedly verified
Automate	Autonomy for proven playbooks	Trust, rollback/verification evidence, policy approval
Fleet	Multi-team / multi-environment governance	Enterprise procurement, RBAC, support, multi-org needs

Stage	Goal	Exit trigger
Transform	Address root infrastructure causes	Recurring architectural faults create engineering case
Retain	Managed infrastructure / continuous optimization	Ongoing operational ownership

7. Business discovery

- How many incidents per month require a human to execute a known fix?
- How many people become involved, for how long, and what business activity is affected?
- What blocks more automation today: trust, authority, security, change control, or verification?
- Which three repetitive incident classes would be safest to prove first?
- What changed recently that makes this more urgent now?

8. Commercial readiness

Validate internally	Reason
Independent production case studies	Trust in autonomous remediation
Security attestations and data-flow documentation	Enterprise procurement
Credential model and failure modes	Architecture and security review
Production coverage by incident class	Avoid overusing replay benchmark
Support / SLA / Fleet governance	Enterprise readiness
Roadmap vs current feature set	Prevent selling future capability

9. Market opportunity context

- AIOps is a growing category, but category definitions vary across analysts. The stronger commercial story is the operational gap between detection and verified remediation.
- Priority accounts: cloud-native SaaS, FinTech, HealthTech, infrastructure-heavy digital businesses and MSPs with 24/7 operations.
- US and UAE are Priority 1 markets for the broader architecture; UK and India are Priority 2, with account selection driven by cloud/SRE triggers rather than geography alone.

10. Real target-account intelligence examples

These are research targets based on current public signals. They are not confirmed opportunities or evidence of buying intent.

Airwallex, UK / EMEA

Fact: In March 2026 Airwallex announced more than \$1 billion of EMEA investment, approximately 100 engineering hires in London, and continued rapid regional transaction growth.

Commercial hypothesis: Platform and engineering expansion can increase change rate, infrastructure footprint and on-call complexity, but the actual incident burden must be validated.

Discovery question: As the EMEA platform footprint and engineering team scale, which recurring production incidents still require manual engineer intervention?

Core42 / e& UAE

Fact: In July 2026 e& UAE and Core42 announced Sovereign AI Compute, combining in-country GPU infrastructure, connectivity and implementation services with an explicit focus on resilience of critical AI services.

Commercial hypothesis: Rapid sovereign AI infrastructure scale creates a credible reliability and controlled-remediation conversation, although a technically sophisticated operator may prefer to build internally.

Discovery question: Which recurring infrastructure events across sovereign AI workloads are already known enough to automate, but are still held back by authority, verification or change-control concerns?

Enterprise Monetization & Account Expansion Architecture

Industry Focus

Industry focus: SaaS, FinTech, HealthTech, e-commerce/digital platforms, cloud-native technology companies, regulated infrastructure teams and MSPs. The strongest fit is where cloud/Kubernetes scale, 24/7 production, repetitive incident remediation, multi-environment operations or client-fleet management create operational leverage.

Documented pricing architecture

Pricing / packaging element	What it means commercially
Individual: first 3 nodes free	Technical evaluation / PLG entry.
Team: \$5/node/month or \$4/node/month annual	Flat per-endpoint pricing; unlimited users and nodes within one organization, RBAC, custom playbooks, multi-environment and API access.
Fleet / MSP: custom	Unlimited organizations / multi-tenancy, custom RBAC, extended audit retention, dedicated success engineer and 24/7 support.
Pricing metric	Per endpoint. SentienGuard explicitly does not meter alerts, custom metrics, user seats, API calls, data transfer or playbook executions.

Enterprise entry and proof

Step	Commercial design
1. Technical entry	Engineer/SRE validates on 3 nodes or limited Team deployment without replacing existing observability.
2. Product-qualified proof	Production environment, node growth, repeatable playbook success, multiple users/teams or API/custom-playbook adoption.
3. Team expansion	More endpoints and environments across the same organization.
4. Fleet / enterprise graduation	Multiple organizations/tenants, custom RBAC, extended audit retention, 24/7 support, dedicated success and MSP/fleet governance.
5. Account depth	Recurring root architectural faults can trigger cloud/platform engineering or managed infrastructure; product telemetry should inform the engineering conversation, not force it.

Enterprise value drivers and expansion triggers

Value driver / trigger	Enterprise implication
Endpoint count	Primary documented pricing metric and linear scaling

	variable.
Environment count	Dev/staging/prod all count; adoption across environments increases breadth.
Organizations / tenants	Multi-tenancy is a decisive Fleet/MSP graduation trigger.
Governance	Custom RBAC and authority policies support enterprise deployment.
Audit retention / compliance evidence	Extended retention is a documented Fleet value driver.
Support	24/7 support and dedicated success are Fleet enterprise requirements.
PQA trigger	Node growth, production use, multiple teams, custom playbooks/API, security/procurement questions.
CEO Input Required	Fleet ACV bands, volume discounts, support economics, renewal/expansion and managed-services attachment.

Enterprise sales rule

Do not lead with the lowest public price. Lead with the business problem, prove value in a contained scope, then expand only when customer scale, complexity, governance, integration or operational ownership creates a legitimate enterprise requirement.

Do not invent ACV, gross margin, discount bands or large-contract claims. Where these are not documented, mark them Executive Input Required.

Updated pricing source: SentienGuard official Pricing page, accessed Aug 2026 - first 3 nodes free; Team \$5 monthly / \$4 annual per node; Fleet/MSP custom; no per-alert/API/metric/playbook metering.

Source notes

[1] SentienGuard official homepage, accessed Aug 2026.

[2] SentienGuard official Pricing page, accessed Aug 2026.

[3] SentienGuard official Automated Remediation page, accessed Aug 2026.

[4] SentienGuard official Infrastructure Automation page, accessed Aug 2026.

[5] SentienGuard official Self-Healing Engine page, accessed Aug 2026.

Public performance metrics are vendor-reported and must retain their stated validation context.