

THE ALGORITHM

SentienGuard GTM Playbook

AI Ops category, autonomy assessment, controlled-remediation motion and enterprise conversion plan.

GTM INTELLIGENCE LIBRARY

Confidential strategic working document. Verified company information, proposed GTM strategy, personal experience and illustrative planning assumptions should be distinguished when used externally.

1. Product understanding

SentienGuard should be positioned as an autonomous remediation layer that works with the customer's existing observability/infrastructure stack. The commercial problem is not lack of alerts; it is the engineering time, MTTR and risk created after the alert fires.

Enter in observation mode, quantify repetitive incidents and engineering toil, automate only controlled/repeatable remediation, verify outcomes, then expand endpoints/environments and attach managed infrastructure.

2. Commercial problem

- Alert fatigue and noisy on-call rotations.
- Repetitive incidents with known remediation.
- High MTTR and escalation burden.
- Runbooks that depend on individual engineers.
- Tool fragmentation across monitoring, incident and infrastructure platforms.
- Need for traceable, controlled automation.

3. Ideal Customer Profile

Priority ICP	Why it matters
Cloud-native SaaS	High uptime expectations, DevOps/SRE maturity and repetitive incident patterns.
FinTech / regulated tech	Reliability plus auditability creates strong business case.
HealthTech	Availability, compliance and lean platform teams.
MSPs	Fleet economics and repeated operational patterns across customers.
Enterprise platform teams	Large estates where toil and MTTR create measurable cost.

4. Buying committee

Persona	Primary concern
CTO / CIO	Reliability, engineering leverage, platform risk
VP Engineering	Productivity, operational burden and delivery capacity
Head of SRE / DevOps	MTTR, alert fatigue, on-call toil and runbook execution
Platform Engineering	Standardization and safe automation
Security / Compliance	Controls, auditability and boundaries

5. Trigger events

- Kubernetes/cloud expansion.

- SRE/DevOps hiring.
- High-profile outage or reliability review.
- Cloud-cost / engineering-productivity initiative.
- Observability-stack investment.
- Compliance/audit requirement.
- MSP customer growth.

6. Positioning and message architecture

Observability tells the team what happened. SentienGuard should be sold as the controlled remediation layer that reduces repetitive operational work after detection, without requiring the customer to replace the tools already trusted for monitoring.

7. Competitive frame

Competitor / category	How to position
Datadog / New Relic / Dynatrace	Complement rather than replace; consume signals and automate response.
PagerDuty / incident tools	Differentiate around execution/remediation after incident routing.
ServiceNow ITOM / AIOps	Compete on deployment simplicity, vendor neutrality and focused remediation where relevant.
Runbook automation	Differentiate with diagnosis, verification, auditability and autonomous progression.

8. Market strategy

Market	Priority	Initial play
United States	P1	SaaS, FinTech, HealthTech, cloud-native enterprise and MSPs
United Kingdom	P1/P2	FinTech, SaaS and regulated infrastructure
UAE/Gulf	P2	Enterprise cloud, government platforms and managed-services opportunities
India	P2/P3	SaaS and MSP ecosystem; selective enterprise accounts

9. Account-selection model

Score every account on ICP fit, active trigger, identifiable pain, access to buying roles and account-expansion potential. Prioritize a small number of high-fit accounts before scaling outbound volume.

Factor	Weight	What good looks like
ICP fit	25%	Strong industry, size and operating fit
Trigger	25%	Visible reason to act now
Pain	20%	Measurable operational or financial problem

Factor	Weight	What good looks like
Access	15%	Reachable buying committee
Expansion	15%	Potential for wider product/services footprint

10. Sales motion

- Identify accounts with cloud/Kubernetes/SRE signals.
- Open with Infrastructure Autonomy Assessment, not “AI Ops” jargon.
- Quantify repetitive incidents, MTTR and engineering effort.
- Start read-only / observation where trust is the barrier.
- Move to controlled remediation with approval.
- Expand automation, endpoints/environments and managed infrastructure.

11. Discovery framework

- How many incidents and alerts occur each month?
- Which incident types repeat?
- What percentage already has a known runbook?
- What is the current MTTR?
- How much on-call and senior-engineer time is consumed?
- What monitoring/incident/infrastructure tools are already used?
- Which production actions can be automated safely?
- What compliance/audit controls govern changes?

12. Qualification

- Repeatable incident classes.
- Technical owner willing to test controlled automation.
- Measurable baseline (MTTR/toil/incidents).
- Supported stack / feasible integration.
- Defined trust and approval boundaries.

13. Demo / workshop / pilot

- Start with observation and incident replay where possible.
- Show one repeatable incident end to end: detect/diagnose/remediate/verify/audit.
- Agree explicit actions that require approval versus autonomy.
- Measure before/after MTTR and engineering touch time.
- Expand only after reliability is demonstrated.

14. Commercial model and business case

- Build ROI around engineer hours, incident recovery time, service risk and tooling/operational overhead.
- Use endpoint/node economics only with current pricing verification.
- For enterprise accounts, package technical validation and governance into the commercial process.

15. Objection handling

Objection	Response direction
"We already use Datadog/Dynatrace."	SentienGuard is not positioned as a replacement; it acts on the signals those systems produce.
"We will not let AI touch production."	Start read-only, then approval mode, and automate only approved, bounded runbooks.
"Our incidents are too unique."	Identify the repetitive subset rather than trying to automate every incident.
"We already have runbooks."	The value is consistent execution, verification and audit, not just documentation.

16. Closing strategy

- Confirm the quantified problem and success criteria before commercial negotiation.
- Map procurement, security, legal and executive approval early.
- Reduce decision risk with a clear first phase, milestones and governance.
- Before signing, define the likely expansion path if the first phase succeeds.

17. Partnerships and channels

- Cloud providers and cloud consultancies.
- MSPs / managed infrastructure providers.
- SRE/DevOps communities and platform-engineering partners.
- Compliance/security advisors where auditability is a key wedge.

18. Expansion and cross-sell

- More endpoints and environments.
- MSP/Fleet deployments.
- Managed infrastructure.
- Cloud/platform engineering.
- Modernization and reliability engineering.
- Compliance and technical-support services.

19. 30 / 60 / 90 execution

Period	Focus	Outputs
['0-30 days', 'Validate technical ICP', 'Audit current users/pipeline, choose target stack/size, define autonomy assessment and baseline metrics']	['31-60 days', 'Activate', 'Named-account outreach, technical assessments, trial/observation deployments, MSP conversations']	['61-90 days', 'Prove', 'Controlled-remediation pilots, quantified before/after results, enterprise proposals and expansion paths']

20. KPIs

- Assessments.
- Trial/observation activations.
- Time to first meaningful incident.
- Trial-to-pilot conversion.
- Pilot-to-paid conversion.
- Endpoints/environments expanded.
- Managed-services attach rate.
- Sales cycle by ICP.

21. Relevant commercial experience

- AI solution-selling experience supports technical discovery and business-outcome positioning.
- IT services background supports cloud/infrastructure cross-sell.
- Account-management orientation supports moving from tool adoption to strategic managed relationship.

22. Revenue planning principle

Start with bottom-up pipeline math: target accounts → conversations → qualified opportunities → pilots/proposals → wins → expansion. Do not commit to revenue without validating actual pricing, baseline conversion rates and sales capacity.

Planning assumptions are illustrative only and must be validated against The Algorithm's actual pricing, pipeline, customer base, margins and conversion data before being used as forecasts or commitments.